

# DEEP LEARNING

## CLASS OVERVIEW

*This course covers the fundamentals of deep neural networks at the graduate level. We introduce multi-layer perceptrons, back-propagation, and automatic differentiation. We will also discuss Convolutional Neural Networks, Recurrent Neural Networks, Transformers, and advanced topics in deep learning. The course will be a combination of lectures, presentations, and machine learning competitions.*

- **Lecturer:** Qi (Rose) Yu ([roseyu@ucsd.edu](mailto:roseyu@ucsd.edu))
- **TA:**
  - Rui (Ray) Wang ([ruw020@ucsd.edu](mailto:ruw020@ucsd.edu))
  - Abhishek Tanpure ([atanpure@ucsd.edu](mailto:atanpure@ucsd.edu))
  - Panini Bhamidipati ([abhamidipati@ucsd.edu](mailto:abhamidipati@ucsd.edu))
- **Lecture:** 3:30 pm - 4:50 pm PT, Tuesday, Thursday, CENTR 212
- **Discussion:** 12:00 pm - 12:50 pm PT, Friday, CSB 002
- **Office Hour:**
  - Rose Yu | 5:00 pm - 6:00 pm | Monday | EBU3B 3208
  - Rui (Ray) Wang | 9:00 am - 10:00 am | Friday | B240A - CSE Basement
  - Abhishek Tanpure | 5:00 pm - 6:00 pm | Tuesday | B270A - CSE Basement
  - Panini Bhamidipati | 5:00 pm - 6:00 pm | Thursday | B270A - CSE Basement
- **Canvas:** <https://canvas.ucsd.edu/courses/45077>
- **Piazza:** <https://piazza.com/ucsd/spring2023/cse251b>



## ***SYLLABUS***

<i>Week 1 (April 3rd)</i>	<i>Introduction and Background</i>	<i>HW 1 release</i>
<i>Week 2 (Apr 10th)</i>	<i>Multi-layer perceptron</i>	
<i>Week 3 (Apr 17th)</i>	<i>Automatic Differentiation</i>	<i>HW 2 release</i>
<i>Week 4 (Apr 24th)</i>	<i>Convolutional neural network</i>	
<i>Week 5 (May 1st)</i>	<i>Recurrent neural network</i>	<i>HW3 release</i>
<i>Week 6 (May 8th)</i>	<i>Mid-term week</i>	
<i>Week 7 (May 15th)</i>	<i>Deep learning implementation</i>	<i>HW 4 release</i>
<i>Week 8 (May 22nd)</i>	<i>Attention and Transformer</i>	<i>Milestone report due</i>
<i>Week 9 (May 29th)</i>	<i>Graph neural network</i>	

# LECTURES

## Lecture Slides

### Latex Template

- [Homework Template](#)
- [Presentation Template](#)
- [Project Template](#)

# CLASS ASSESSMENT

- 40 % homework (10% x 4)
- 35 % Kaggle competition
  - 5 % milestone report
  - 10 % final report
  - 10 % final presentation
  - 10 % competition ranking
- 25 % Mid-term exam

# RESOURCES

## Kaggle Competition

- [Competition Page](#)

## Reading Materials

- [1] [Machine Learning Crash Course](#)
- [2] [Deep Learning Book](#)
- [3] [Learning from Data](#)

## Other Tutorials

- [Deep Generative Models.](#)
- [Generative Adversarial Networks.](#)
- [Dive Deep Into Deep Learning](#)

# FAQ

## ***Q: What are the pre-requisites?***

- *(MATH 31BH or MATH 20C) and (ECON 120A or ECE 109 or CSE 103 or MATH 181A or MATH 183, MATH 170A);*
- *Proficiency in Python.*

## ***Q: Can first year undergraduates take this course?***

- *Restricted to students with sophomore, junior, or senior standing within the CS25, CS26, CS27, CS28, EC26, and DS25 majors.*
- *All other students will be allowed as space permits.*



## ***About me***

*My Chinese name is Qi Yu. That is also the instructor's name in the registrar's office. I publish under the name Rose Yu. You can learn more about my research at [my personal website](#).*

*Rose Yu   [roseyu@eng.ucsd.edu](mailto:roseyu@eng.ucsd.edu)*